

# CLASSIFIER - JSMA - FGSM

November 15, 2020

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import keras

session = tf.Session()
keras.backend.set_session(session)
```

Using TensorFlow backend.

```
[2]: x_train = np.load("./DATA/X_train.npy", allow_pickle=True)
x_test = np.load("./DATA/X_test.npy", allow_pickle=True)
y_train = np.load("./DATA/y_train.npy", allow_pickle=True)
y_test = np.load("./DATA/y_test.npy", allow_pickle=True)
```

```
[3]: y_test.shape
```

```
[3]: (5000, 1)
```

```
[4]: print ("Training Examples: %d" % len(x_train))
print ("Test Examples: %d" % len(x_test))
```

Training Examples: 10000

Test Examples: 5000

```
[5]: n_classes = 2
inds=np.array([y_train==i for i in range(n_classes)])
f,ax=plt.subplots(2,5,figsize=(10,5))
ax=ax.flatten()
for i in range(n_classes):
    ax[i].imshow(x_train[np.argmax(inds[i])].reshape(48,48))
    ax[i].set_title(str(i))
plt.show()
```



---

---

---

---

---

---

## 1 Classifier

We build the classifier that will be used to evaluate the testing and attacks

<Describe the classifier> - BALAJI Specialization. We used convolution - remember

```
[6]: from keras import models
from keras import layers

network = models.Sequential()
network.add(layers.Reshape((48*48,),input_shape=(48,48,)))
network.add(layers.Dense(512, activation='relu', input_shape=(48 * 48,)))
network.add(layers.Dense(2, activation='softmax'))

network.compile(optimizer='rmsprop',
                loss='categorical_crossentropy',
                metrics=['accuracy'])

network.summary()
```

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\tensorflow\_core\python\ops\resource\_variable\_ops.py:1630: calling BaseResourceVariable.\_\_init\_\_ (from tensorflow.python.ops.resource\_variable\_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass \*\_constraint arguments to layers.

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
reshape_1 (Reshape)	(None, 2304)	0
dense_1 (Dense)	(None, 512)	1180160
dense_2 (Dense)	(None, 2)	1026

Total params: 1,181,186  
Trainable params: 1,181,186

Non-trainable params: 0

-----

```
[7]: from keras.utils import to_categorical #this just converts the labels to  
      ↪ one-hot class  
      train_labels = to_categorical(y_train)  
      test_labels = to_categorical(y_test)
```

```
[8]: from keras.callbacks import ModelCheckpoint  
  
      h=network.fit(x_train,  
                    train_labels,  
                    epochs=5,  
                    batch_size=256,  
                    shuffle=True,  
                    callbacks=[ModelCheckpoint('image_space.h5', save_best_only=True)])
```

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\keras\backend\tensorflow\_backend.py:422: The name tf.global\_variables is deprecated. Please use tf.compat.v1.global\_variables instead.

Epoch 1/5  
10000/10000 [=====] - 1s 72us/step - loss: 0.2328 -  
accuracy: 0.9171

Epoch 2/5  
3072/10000 [=====>...] - ETA: 0s - loss: 0.1615 -  
accuracy: 0.9443

C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\keras\callbacks\callbacks.py:707: RuntimeWarning: Can save best model only with val\_loss available, skipping.  
'skipping.' % (self.monitor), RuntimeWarning)

10000/10000 [=====] - 0s 39us/step - loss: 0.1519 -  
accuracy: 0.9493


Epoch 3/5  
10000/10000 [=====] - 0s 41us/step - loss: 0.1161 -  
accuracy: 0.9601

Epoch 4/5  
10000/10000 [=====] - 0s 39us/step - loss: 0.0982 -  
accuracy: 0.9641

Epoch 5/5  
10000/10000 [=====] - 0s 46us/step - loss: 0.0828 -  
accuracy: 0.9720

```
[9]: # summarize history for accuracy  
      plt.plot(h.history['accuracy'])  
      plt.title('model accuracy')
```

```
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train'], loc='upper left')
plt.show()
```



```
[10]: network.evaluate(x_train, train_labels)
```

10000/10000 [=====] - 1s 64us/step

```
[10]: [0.08402268831133843, 0.9724000096321106]
```

```
[11]: score, acc = network.evaluate(x_test, test_labels,)
print ("Test Accuracy: %.5f" % acc)
```

5000/5000 [=====] - 0s 64us/step

Test Accuracy: 0.94760

## 1.1 SAVE CLASSIFIER

```
[12]: network.save("CLASSIFIER.h5py")
```

## 1.2 Find all MAL and BEN Rows

```
[13]: MAL = []
      BEN = []
      c = cc = 0
      for i in range(len(test_labels[:,1])):
          if test_labels[i][1] == 0:
              BEN.append(i)
          else:
              MAL.append(i)
```

[illegible]

## 2 FGSM

```
[14]: from cleverhans.utils_keras import KerasModelWrapper
      wrap = KerasModelWrapper(network)
```

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\cleverhans\utils\_tf.py:341: The name tf.GraphKeys is deprecated. Please use tf.compat.v1.GraphKeys instead.

```
[15]: from cleverhans.attacks import FastGradientMethod
      fgsm = FastGradientMethod(wrap, sess=session)
```

```
[16]: x = tf.placeholder(tf.float32, shape=(None, 2304))
      y = tf.placeholder(tf.float32, shape=(None, 2))
```

```
[17]: fgsm_rate = 0.08
      fgsm_params = {'eps': fgsm_rate, 'clip_min': 0., 'clip_max': 1.}
```

```
[18]: test_images_mal = x_test[MAL].copy()
```

```
[19]: adv_x = fgsm.generate_np(test_images_mal, **fgsm_params)
```

[INFO 2020-11-15 17:08:00,706 cleverhans] Constructing new graph for attack FastGradientMethod

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\cleverhans\attacks\\_init\_\_.py:283: to\_float (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use `tf.cast` instead.

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\cleverhans\utils\_tf.py:624: The name tf.assert\_greater\_equal is deprecated. Please use tf.compat.v1.assert\_greater\_equal instead.

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\cleverhans\utils\_tf.py:615: The name tf.assert\_less\_equal is deprecated. Please use tf.compat.v1.assert\_less\_equal instead.

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\cleverhans\compat.py:124: calling softmax\_cross\_entropy\_with\_logits\_v2\_helper (from tensorflow.python.ops.nn\_ops) with dim is deprecated and will be removed in a future version.

Instructions for updating:

dim is deprecated, use axis instead

```
[20]: for i in range(test_images_mal.shape[0]):  
      for j in range(test_images_mal.shape[1]):  
        for k in range(test_images_mal.shape[2]):  
          if test_images_mal[i][j][k]==1:  
            adv_x[i][j][k] = 1
```

```
[21]: network.evaluate(test_images_mal, test_labels[MAL], batch_size=128)
```

1667/1667 [=====] - 0s 76us/step

```
[21]: [0.36963911920374715, 0.9064187407493591]
```

```
[22]: network.evaluate(adv_x, test_labels[MAL], batch_size=128)
```

1667/1667 [=====] - 0s 38us/step

```
[22]: [24.913951785829372, 0.09298140555620193]
```

```
[23]: x_test_after_attack_FGSM = x_test.copy()
```

```
[24]: x_test_after_attack_FGSM[MAL] = adv_x
```

```
[25]: network.evaluate(x_test, test_labels, batch_size=128)
```

5000/5000 [=====] - 0s 43us/step

```
[25]: [0.19247934894561766, 0.9476000070571899]
```

```
[26]: network.evaluate(x_test_after_attack_FGSM, test_labels, batch_size=128)
```

5000/5000 [=====] - 0s 42us/step

```
[26]: [8.375553169500828, 0.6764000058174133]
```

## 2.1 SAVE ALL VARS

- X\_TEST
- Y\_TEST
- X\_TEST\_AFTER\_ATTACK\_FGSM

```
[27]: np.save('./ATTACKS/FGSM/X_TEST_FGSM.npy', x_test)  
      np.save('./ATTACKS/FGSM/Y_TEST_FGSM.npy', y_test)  
      np.save('./ATTACKS/FGSM/X_TEST_ATTACKED_FGSM.npy', x_test_after_attack_FGSM)  
      network.save('./ATTACKS/FGSM/FGSM_CLASSIFIER_USED.h5py')
```

---

---



### 3 LOAD MODEL

```
[29]: network = keras.models.load_model('CLASSIFIER.h5py')
      network.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

reshape_1 (Reshape)	(None, 2304)	0
-----		
dense_1 (Dense)	(None, 512)	1180160
-----		
dense_2 (Dense)	(None, 2)	1026
=====		
Total params: 1,181,186		
Trainable params: 1,181,186		
Non-trainable params: 0		
-----		

## 4 JSMA

```
[30]: from cleverhans.attacks import SaliencyMapMethod
      from cleverhans.utils_keras import KerasModelWrapper
      wrap = KerasModelWrapper(network)
      jsma = SaliencyMapMethod(wrap, sess=session)

[31]: # x = tf.placeholder(tf.float32, shape=(None, 2304))
      # y = tf.placeholder(tf.float32, shape=(None, 2))

[32]: jsma_params = {'theta': 1.,
                    'gamma': 0.1,
                    'clip_min': 0.,
                    'clip_max': 1.,
                    'y_target': None}

[33]: y_test.shape

[33]: (5000, 1)

[34]: from keras.utils import to_categorical #this just converts the labels to
      ↪one-hot class
      test_labels = to_categorical(y_test)

[35]: x_test_afterattack=np.zeros(x_test.shape)

      crafted=[]

[36]: nb_classes = 2

[37]: x_test_mal_noisy = []
      x_test_mal_noisy_idx = []
      for i in range(x_test.shape[0]):
          if i in MAL:
              sample = x_test[i: i + 1]
              one_hot_target = np.zeros((1, nb_classes), dtype=np.float32)
```

```

        one_hot_target[0, 0] = 1
        jsma_params['y_target'] = one_hot_target
        adv_x = jsma.generate_np(sample, **jsma_params)
        x_test_afterattack[i]=adv_x
    else:
        x_test_afterattack[i] = x_test[i]
    if i%100 == 0:
        print("AT : " + str(i))

```

[INFO 2020-11-15 17:08:12,777 cleverhans] Constructing new graph for attack  
SaliencyMapMethod

AT : 0

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\cleverhans\attacks\_tf.py:446: The name tf.mod is deprecated. Please use tf.math.mod instead.

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\cleverhans\attacks\_tf.py:447: The name tf.floordiv is deprecated. Please use tf.math.floordiv instead.

AT : 100  
 AT : 200  
 AT : 300  
 AT : 400  
 AT : 500  
 AT : 600  
 AT : 700  
 AT : 800  
 AT : 900  
 AT : 1000  
 AT : 1100  
 AT : 1200  
 AT : 1300  
 AT : 1400  
 AT : 1500  
 AT : 1600  
 AT : 1700  
 AT : 1800  
 AT : 1900  
 AT : 2000  
 AT : 2100  
 AT : 2200  
 AT : 2300  
 AT : 2400  
 AT : 2500  
 AT : 2600  
 AT : 2700

AT : 2800  
AT : 2900  
AT : 3000  
AT : 3100  
AT : 3200  
AT : 3300  
AT : 3400  
AT : 3500  
AT : 3600  
AT : 3700  
AT : 3800  
AT : 3900  
AT : 4000  
AT : 4100  
AT : 4200  
AT : 4300  
AT : 4400  
AT : 4500  
AT : 4600  
AT : 4700  
AT : 4800  
AT : 4900

```
[38]: # adv_x = jsma.generate_np(x_test[MAL], **jsma_params)
```

```
[39]: x_test_afterattack.shape
```

```
[39]: (5000, 48, 48)
```

```
[40]: network.evaluate(x_test, test_labels, batch_size=128)
```

5000/5000 [=====] - 0s 50us/step

```
[40]: [0.19247934894561766, 0.9476000070571899]
```

```
[41]: network.evaluate(x_test_afterattack, test_labels, batch_size=128)
```

5000/5000 [=====] - 0s 39us/step

```
[41]: [0.8315977962493897, 0.6453999876976013]
```

```
[42]: x_test_after_attack_JSMA = x_test_afterattack.copy()
```

#### 4.1 SAVE ALL VARS

- X\_TEST
- Y\_TEST
- X\_TEST\_AFTER\_ATTACK\_FGSM

```
[43]: np.save('./ATTACKS/JSMA/X_TEST_JSMA.npy',x_test)
      np.save('./ATTACKS/JSMA/Y_TEST_JSMA.npy',y_test)
      np.save('./ATTACKS/JSMA/X_TEST_ATTACKED_JSMA.npy',x_test_after_attack_JSMA)
      network.save('./ATTACKS/JSMA/JSMA_CLASSIFIER_USED.h5py')
```

**5 Pretty straightforward. Talk about Attacks, Drop in acc. Rest dealt in Attack EDA**

```
[ ]:
```