

# APE GAN FGSM PURE MAL

November 14, 2020

```
[2]: print("STARTING")
```

STARTING

```
[3]: import numpy as np
from sklearn.model_selection import train_test_split
import random
import numpy as np
import keras
import tensorflow as tf
from keras.datasets import mnist
from keras.utils import np_utils # utilities for one-hot encoding of ground
    → truth values

import os

import keras
from keras.models import Model, Sequential # basic class for specifying and
    → training a neural network
from keras.layers import Input, Conv2D, Conv2DTranspose, Dense, Activation,
    → Flatten, LeakyReLU, BatchNormalization, ZeroPadding2D, Conv1D
from keras.optimizers import Adam
from keras import backend as K

os.environ["CUDA_VISIBLE_DEVICES"]="1"

%reload_ext autoreload
%autoreload 2
```

Using TensorFlow backend.

```
[4]: X_TEST = np.load('./DATA/X_CLEAN_ONLY_MAL.npy')
Y_TEST = np.load('./DATA/X_LABEL_ONLY_MAL.npy')
coeff = np.load('./DATA/ORIGINAL/coeff_features.npy')
X_N_TEST = np.load('./DATA/X_ADV_ONLY_MAL.npy')
```

# 1 Classifier Import and test

Building a classifier to evaluate the denoising

```
[6]: np.unique(Y_TEST, return_counts=True)
```

```
[6]: (array([1]), array([3799], dtype=int64))
```

```
[7]: modelClassifier = keras.models.load_model('./modelClassifierFGSM.h5')
modelClassifier.summary()
```

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\tensorflow\_core\python\ops\resource\_variable\_ops.py:1630: calling BaseResourceVariable.\_\_init\_\_ (from tensorflow.python.ops.resource\_variable\_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass \*\_constraint arguments to layers.

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\tensorflow\_core\python\ops\nn\_impl.py:183: where (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-packages\keras\backend\tensorflow\_backend.py:422: The name tf.global\_variables is deprecated. Please use tf.compat.v1.global\_variables instead.

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 2948)	8693652
dense_8 (Dense)	(None, 128)	377472
dense_9 (Dense)	(None, 1)	129

Total params: 9,071,253  
Trainable params: 9,071,253  
Non-trainable params: 0

```
[9]: modelClassifier.evaluate(X_TEST, Y_TEST)
```

```
3799/3799 [=====] - 2s 556us/step
```

```
[9]: [0.04282340146973827, 1.0]
```

```
[10]: modelClassifier.evaluate(X_N_TEST, Y_TEST)
```

```
3799/3799 [=====] - 2s 450us/step
```

```
[10]: [299.8835265824091, 0.0]
```

### 1.0.1 Accuracy of classifier

- Clean data = 100%
- FGSM attacked data = 0%

But remeber this is just malware

```
[11]: weightsmodelClassifier, biasesmodelClassifier = modelClassifier.layers[0].  
      ↪get_weights()  
      modelClassifier.layers[0].get_weights()[1].shape
```

```
[11]: (2948,)
```

## 2 Ben and Mal Cols and Pure Malware and Benign Arrays

Pure Malware = all malware columns will be one

Pure Ben = all Ben cols will be one

This gives us the 2 extreme cases to compare with one another during GAN training. All Ben being the most benign and all mal being the worst

```
[12]: mal_col_index = []  
      ben_col_index = []  
      for i in range(len(coeff)):  
          if coeff[i] > 0:  
              mal_col_index.append(i)  
          elif coeff[i] < 0:  
              ben_col_index.append(i)  
          else:  
              print("DANGER")  
      print(len(mal_col_index), len(ben_col_index))
```

```
1513 1435
```

```
[13]: ALL_BEN_APK = np.zeros(X_TEST[0].shape)  
      for i in ben_col_index:  
          ALL_BEN_APK[i] = 1  
      np.unique(ALL_BEN_APK, return_counts=True)
```

```
[13]: (array([0., 1.]), array([1513, 1435], dtype=int64))
```

```
[14]: ALL_MAL_APK = np.zeros(X_TEST[0].shape)
      for i in mal_col_index:
          ALL_MAL_APK[i] = 1
      np.unique(ALL_MAL_APK, return_counts=True)
```

```
[14]: (array([0., 1.]), array([1435, 1513], dtype=int64))
```

### 3 APE GAN DEF

```
[15]: def MANHATTAN(y_true, y_pred):
      return K.sum( K.abs( y_true - y_pred), axis=1, keepdims=True) + 1e-10

      def SRMSE(y_true, y_pred):
          return K.sqrt(K.mean(K.square(y_pred - y_true), axis=-1) + 1e-10)

      def WMOD(y_true, y_pred):
          return K.abs(1 - K.mean(y_true * y_pred))

      def WGAN(y_true, y_pred):
          return K.abs(K.mean(y_true * y_pred))
```

```
[16]: def generator(input_dims):
      model = Sequential()
      model.add(Dense(512, input_shape = input_dims, activation='relu'))
      #model.add(Dense(2048, input_shape = input_dims, activation='relu'))
      #model.add(Dense(1024, activation='relu'))
      #model.add(Dense(512, activation='relu'))
      model.add(Dense(256, activation='relu'))
      model.add(Dense(128, activation='relu'))
      model.add(Dense(64, activation='relu'))
      model.add(Dense(32, activation='relu'))
      model.add(Dense(16, activation='relu'))
      model.add(Dense(8, activation='relu'))
      model.add(Dense(4, activation='relu'))
      model.add(Dense(2, activation='relu'))
      model.add(Dense(1, activation='relu'))
      model.add(Activation('tanh'))
      return model

      def discriminator(input_dims):
          model = Sequential()
          model.add(Dense(512, input_shape = input_dims, activation='relu'))
          #model.add(Dense(2048, input_shape = input_dims, activation='relu'))
          #model.add(Dense(1024, activation='relu'))
          #model.add(Dense(512, activation='relu'))
          model.add(Dense(256, activation='relu'))
```

```

model.add(Dense(128, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(Dense(16, activation='relu'))
model.add(Dense(8, activation='relu'))
model.add(Dense(4, activation='relu'))
model.add(Dense(2, activation='relu'))
model.add(Dense(1, activation='relu'))
model.add(Activation('sigmoid'))
model.compile(optimizer='adam', loss='binary_crossentropy')
return model

def APEGAN(input_dims):
    G = generator(input_dims)
    print("=====\n\nGENERATOR\n\n")
    print(G.summary())
    print("=====\n\n")

    D = discriminator(input_dims)
    print("=====\n\nDISCRIMINATOR\n\n")
    print(D.summary())
    print("=====\n\n")

    ipt = Input(input_dims)
    print("=====\n\nINPUT TENSOR\n\n")
    print(ipt)
    print("=====\n\n")

    purified = G(ipt)
    print("=====\n\nPURIFIED TENSOR\n\n")
    print(purified)
    print("=====\n\n")

    D.trainable = False

    judge = D(purified)
    print("=====\n\nJUDGE TENSOR\n\n")
    print(judge)
    print("=====\n\n")

    GAN = Model(ipt, [judge, purified])
    print("=====\n\nGAN BASIC\n\n")
    print(GAN.summary())
    print("=====\n\n")

    GAN.compile(optimizer='adam',
                loss=['binary_crossentropy', WGAN],

```

```

        loss_weights=[0.02, 0.9])

    print("=====\n\nGAN AFTER COMPILE\n\n")
    print(GAN.summary())
    print("=====\n\n")

    return GAN,G,D

```

## 4 APE GAN RUN

```

[17]: epochs=10 # original 500
      batch_size=128

```

```

N = X_TEST.shape[0]
x_clean = X_TEST.copy()
x_adv = X_N_TEST.copy()
x_label = Y_TEST.copy()

```

```

[18]: GAN, G, D = APEGAN([1,X_TEST.shape[1],1])
      # GAN,G,D = APEGAN([X_TEST.shape[1]])

```

=====

GENERATOR

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 1, 2948, 512)	1024
dense_2 (Dense)	(None, 1, 2948, 256)	131328
dense_3 (Dense)	(None, 1, 2948, 128)	32896
dense_4 (Dense)	(None, 1, 2948, 64)	8256
dense_5 (Dense)	(None, 1, 2948, 32)	2080
dense_6 (Dense)	(None, 1, 2948, 16)	528
dense_7 (Dense)	(None, 1, 2948, 8)	136
dense_8 (Dense)	(None, 1, 2948, 4)	36

dense_9 (Dense)	(None, 1, 2948, 2)	10
-----		
dense_10 (Dense)	(None, 1, 2948, 1)	3
-----		
activation_1 (Activation)	(None, 1, 2948, 1)	0
=====		
Total params: 176,297		
Trainable params: 176,297		
Non-trainable params: 0		
-----		
None		
=====		

## DISCRIMINATOR

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
dense_11 (Dense)	(None, 1, 2948, 512)	1024
-----		
dense_12 (Dense)	(None, 1, 2948, 256)	131328
-----		
dense_13 (Dense)	(None, 1, 2948, 128)	32896
-----		
dense_14 (Dense)	(None, 1, 2948, 64)	8256
-----		
dense_15 (Dense)	(None, 1, 2948, 32)	2080
-----		
dense_16 (Dense)	(None, 1, 2948, 16)	528
-----		
dense_17 (Dense)	(None, 1, 2948, 8)	136
-----		
dense_18 (Dense)	(None, 1, 2948, 4)	36
-----		
dense_19 (Dense)	(None, 1, 2948, 2)	10
-----		
dense_20 (Dense)	(None, 1, 2948, 1)	3
-----		
activation_2 (Activation)	(None, 1, 2948, 1)	0
=====		
Total params: 176,297		
Trainable params: 176,297		
Non-trainable params: 0		

-----  
None  
=====

=====

INPUT TENSOR

Tensor("input\_1:0", shape=(?, 1, 2948, 1), dtype=float32)  
=====

=====

PURIFIED TENSOR

Tensor("sequential\_1/activation\_1/Tanh:0", shape=(?, 1, 2948, 1), dtype=float32)  
=====

=====

JUDGE TENSOR

Tensor("sequential\_2/activation\_2/Sigmoid:0", shape=(?, 1, 2948, 1),  
dtype=float32)  
=====

=====

GAN BASIC

Model: "model\_1"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 1, 2948, 1)	0
sequential_1 (Sequential)	(None, 1, 2948, 1)	176297
sequential_2 (Sequential)	(None, 1, 2948, 1)	176297



Total params: 352,594  
 Trainable params: 176,297  
 Non-trainable params: 176,297

-----  
 None

=====

GAN AFTER COMPILE

Model: "model\_1"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 1, 2948, 1)	0
-----		
sequential_1 (Sequential)	(None, 1, 2948, 1)	176297
-----		
sequential_2 (Sequential)	(None, 1, 2948, 1)	176297
=====		

Total params: 352,594  
 Trainable params: 176,297  
 Non-trainable params: 176,297

-----  
 None

```
[19]: scalarloss = [0,0,0]
      LossHistory = [0]*epochs
      for cur_epoch in range(epochs):

          print("\n\nEPOCH",cur_epoch,"\n")

          idx = np.random.randint(0, N*4//5, size=batch_size)
          x_clean_batch = x_clean[idx,:].reshape(-1,1,2948,1)
          x_adv_batch = x_adv[idx,:].reshape(-1,1,2948,1)

          ALL_MAL_APK_BATCH = np.tile(ALL_MAL_APK,(batch_size,1))
          ALL_BEN_APK_BATCH = np.tile(ALL_BEN_APK,(batch_size,1))

          print("\n=====\nSHAPE of XCLEAN")
          print(x_clean_batch.shape)
```

```

    scalarloss[0] = D.train_on_batch(x_clean_batch, ALL_MAL_APK_BATCH.
↪reshape(-1,1,2948,1))/2

    print("\n====\nNP ONES TRAIN ON BATCH DISCRIMINATOR")
    print("1 "+str(scalarloss))

    scalarloss[0] += D.train_on_batch(x_adv_batch, ALL_BEN_APK_BATCH.
↪reshape(-1,1,2948,1))/2

    print("\n====\nNP ZEROS TRAIN ON BATCH DISCRIMINATOR")
    print("2 "+str(scalarloss))

    GAN.train_on_batch(x_adv_batch, [ ALL_MAL_APK_BATCH.reshape(-1,1,2948,1),
↪x_clean_batch])

    scalarloss[1:] = GAN.train_on_batch(x_adv_batch, [ ALL_MAL_APK_BATCH.
↪reshape(-1,1,2948,1), x_clean_batch])[1:]

#     LossHistory.append(scalarloss)
    LossHistory[cur_epoch] = scalarloss
    print("\n====\nLOSS HISTORY")
    print(LossHistory)

    print("\n====\nTRAIN ON BATCH GAN")
    print("Epoch number:",cur_epoch,"; Loss",scalarloss)
    print("\n EPOCHING \n\n\n\n\n")

```

EPOCH 0

=====

SHAPE of XCLEAN  
(128, 1, 2948, 1)

C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-  
packages\keras\engine\training.py:297: UserWarning: Discrepancy between  
trainable weights and collected trainable weights, did you set `model.trainable`  
without calling `model.compile` after ?  
'Discrepancy between trainable weights and collected trainable'

=====

NP ONES TRAIN ON BATCH DISCRIMINATOR  
1 [0.34657296538352966, 0, 0]

C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-

```
packages\keras\engine\training.py:297: UserWarning: Discrepancy between
trainable weights and collected trainable weights, did you set `model.trainable`
without calling `model.compile` after ?
```

```
'Discrepancy between trainable weights and collected trainable'
```

```
=====
```

```
NP ZEROS TRAIN ON BATCH DISCRIMINATOR
```

```
2 [0.6931459307670593, 0, 0]
```

```
=====
```

```
LOSS HISTORY
```

```
[[0.6931459307670593, 0.69314593, 0.0], 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

```
=====
```

```
TRAIN ON BATCH GAN
```

```
Epoch number: 0 ; Loss [0.6931459307670593, 0.69314593, 0.0]
```

```
EPOCHING
```

```
EPOCH 1
```

```
=====
```

```
SHAPE of XCLEAN
```

```
(128, 1, 2948, 1)
```

```
C:\Users\Pitch\.conda\envs\tf1-gpu\lib\site-
```

```
packages\keras\engine\training.py:297: UserWarning: Discrepancy between
trainable weights and collected trainable weights, did you set `model.trainable`
without calling `model.compile` after ?
```

```
'Discrepancy between trainable weights and collected trainable'
```

```
=====
```

```
NP ONES TRAIN ON BATCH DISCRIMINATOR
```

```
1 [0.34657296538352966, 0.69314593, 0.0]
```

```
=====
```

```
NP ZEROS TRAIN ON BATCH DISCRIMINATOR
```

```
2 [0.6931459307670593, 0.69314593, 0.0]
```

```
=====
```

```
LOSS HISTORY
[[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],
0, 0, 0, 0, 0, 0, 0, 0]
```

```
=====
```

```
TRAIN ON BATCH GAN
```

```
Epoch number: 1 ; Loss [0.6931459307670593, 0.69314593, 0.0]
```

```
EPOCHING
```

```
EPOCH 2
```

```
=====
```

```
SHAPE of XCLEAN
(128, 1, 2948, 1)
```

```
=====
```

```
NP ONES TRAIN ON BATCH DISCRIMINATOR
```

```
1 [0.34657296538352966, 0.69314593, 0.0]
```

```
=====
```

```
NP ZEROS TRAIN ON BATCH DISCRIMINATOR
```

```
2 [0.6931459307670593, 0.69314593, 0.0]
```

```
=====
```

```
LOSS HISTORY
```

```
[[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],
[0.6931459307670593, 0.69314593, 0.0], 0, 0, 0, 0, 0, 0, 0, 0]
```

```
=====
```

```
TRAIN ON BATCH GAN
```

```
Epoch number: 2 ; Loss [0.6931459307670593, 0.69314593, 0.0]
```

```
EPOCHING
```

EPOCH 3

=====

SHAPE of XCLEAN  
(128, 1, 2948, 1)

=====

NP ONES TRAIN ON BATCH DISCRIMINATOR  
1 [0.34657296538352966, 0.69314593, 0.0]

=====

NP ZEROS TRAIN ON BATCH DISCRIMINATOR  
2 [0.6931459307670593, 0.69314593, 0.0]

=====

LOSS HISTORY  
[[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0], 0,  
0, 0, 0, 0, 0]

=====

TRAIN ON BATCH GAN  
Epoch number: 3 ; Loss [0.6931459307670593, 0.69314593, 0.0]

EPOCHING

EPOCH 4

=====

SHAPE of XCLEAN  
(128, 1, 2948, 1)

=====

NP ONES TRAIN ON BATCH DISCRIMINATOR  
1 [0.34657296538352966, 0.69314593, 0.0]

=====

NP ZEROS TRAIN ON BATCH DISCRIMINATOR  
2 [0.6931459307670593, 0.69314593, 0.0]

=====

LOSS HISTORY

[[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], 0, 0, 0, 0, 0]

=====

TRAIN ON BATCH GAN

Epoch number: 4 ; Loss [0.6931459307670593, 0.69314593, 0.0]

EPOCHING

EPOCH 5

=====

SHAPE of XCLEAN

(128, 1, 2948, 1)

=====

NP ONES TRAIN ON BATCH DISCRIMINATOR

1 [0.34657296538352966, 0.69314593, 0.0]

=====

NP ZEROS TRAIN ON BATCH DISCRIMINATOR

2 [0.6931459307670593, 0.69314593, 0.0]

=====

LOSS HISTORY

[[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0], 0,  
0, 0, 0]

=====

TRAIN ON BATCH GAN

Epoch number: 5 ; Loss [0.6931459307670593, 0.69314593, 0.0]

EPOCHING

EPOCH 6

=====

SHAPE of XCLEAN  
(128, 1, 2948, 1)

=====

NP ONES TRAIN ON BATCH DISCRIMINATOR  
1 [0.34657296538352966, 0.69314593, 0.0]

=====

NP ZEROS TRAIN ON BATCH DISCRIMINATOR  
2 [0.6931459307670593, 0.69314593, 0.0]

=====

LOSS HISTORY  
[[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], 0, 0, 0]

=====

TRAIN ON BATCH GAN  
Epoch number: 6 ; Loss [0.6931459307670593, 0.69314593, 0.0]

EPOCHING

EPOCH 7

=====

SHAPE of XCLEAN  
(128, 1, 2948, 1)

=====

NP ONES TRAIN ON BATCH DISCRIMINATOR  
1 [0.34657296538352966, 0.69314593, 0.0]

=====

NP ZEROS TRAIN ON BATCH DISCRIMINATOR  
2 [0.6931459307670593, 0.69314593, 0.0]

=====

LOSS HISTORY  
[[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0], 0,  
0]

=====

TRAIN ON BATCH GAN  
Epoch number: 7 ; Loss [0.6931459307670593, 0.69314593, 0.0]

EPOCHING

EPOCH 8

=====

SHAPE of XCLEAN  
(128, 1, 2948, 1)

=====

NP ONES TRAIN ON BATCH DISCRIMINATOR  
1 [0.34657296538352966, 0.69314593, 0.0]

=====

NP ZEROS TRAIN ON BATCH DISCRIMINATOR  
2 [0.6931459307670593, 0.69314593, 0.0]

=====

LOSS HISTORY  
[[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], 0]



=====

TRAIN ON BATCH GAN

Epoch number: 8 ; Loss [0.6931459307670593, 0.69314593, 0.0]

EPOCHING

EPOCH 9

=====

SHAPE of XCLEAN

(128, 1, 2948, 1)

=====

NP ONES TRAIN ON BATCH DISCRIMINATOR

1 [0.34657296538352966, 0.69314593, 0.0]

=====

NP ZEROS TRAIN ON BATCH DISCRIMINATOR

2 [0.6931459307670593, 0.69314593, 0.0]

=====

LOSS HISTORY

[[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0],  
[0.6931459307670593, 0.69314593, 0.0], [0.6931459307670593, 0.69314593, 0.0]]

=====

TRAIN ON BATCH GAN

Epoch number: 9 ; Loss [0.6931459307670593, 0.69314593, 0.0]

EPOCHING

## 5 Evaluation

We use the Generator of the trained APEGAN to purify data samples and then test it out with the trained classifier

```
[20]: F = keras.models.load_model('./modelClassifierFGSM.h5')
print("=====\n\nKERAS LOAD MODEL\n\n")
print(F.summary())
print("=====\n\n")

clean = X_TEST.copy().reshape(-1,1,X_TEST.shape[1],1)
adv = X_N_TEST.copy().reshape(-1,1,X_TEST.shape[1],1)
label = Y_TEST.copy()

print("=====\n\nCLEAN\n\n")
print(clean)
print("=====\n\n")

print("=====\n\nADV\n\n")
print(adv)
print("=====\n\n")

purified = G.predict(adv)
print("=====\n\nG Predict ADV - PURIFIED\n\n")
print(purified)
print("=====\n\n")
```

=====

KERAS LOAD MODEL

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 2948)	8693652
dense_8 (Dense)	(None, 128)	377472
dense_9 (Dense)	(None, 1)	129

=====

Total params: 9,071,253  
Trainable params: 9,071,253  
Non-trainable params: 0

-----  
None  
=====

=====

CLEAN

```
[[[0.]  
  [0.]  
  [0.]  
  ...  
  [0.]  
  [0.]  
  [0.]]]
```

```
[[[0.]  
  [0.]  
  [0.]  
  ...  
  [0.]  
  [0.]  
  [0.]]]
```

```
[[[0.]  
  [1.]  
  [0.]  
  ...  
  [0.]  
  [0.]  
  [0.]]]
```

...

```
[[[0.]  
  [0.]  
  [0.]  
  ...  
  [0.]  
  [0.]  
  [0.]]]
```

```
[[[0.]  
  [0.]
```

[0.]  
...  
[0.]  
[0.]  
[0.]]]

[[[0.]  
[0.]  
[0.]  
...  
[0.]  
[0.]  
[0.]]]]

=====

=====

ADV

[[[0.]  
[0.]  
[0.]  
...  
[0.]  
[1.]  
[1.]]]

[[[1.]  
[1.]  
[0.]  
...  
[1.]  
[1.]  
[0.]]]

[[[0.]  
[1.]  
[0.]  
...  
[1.]  
[1.]  
[1.]]]

...

```
[[[0.]
  [0.]
  [0.]
  ...
  [0.]
  [1.]
  [1.]]]
```

```
[[[0.]
  [0.]
  [0.]
  ...
  [0.]
  [1.]
  [1.]]]
```

```
[[[1.]
  [0.]
  [0.]
  ...
  [1.]
  [1.]
  [1.]]]
```

=====

=====

G Predict ADV - PURIFIED

```
[[[[0.]
  [0.]
  [0.]
  ...
  [0.]
  [0.]
  [0.]]]
```

```
[[[0.]
  [0.]
```

```

[0.]
...
[0.]
[0.]
[0.]]]

[[[0.]
  [0.]
  [0.]
...
  [0.]
  [0.]
  [0.]]]

...

[[[0.]
  [0.]
  [0.]
...
  [0.]
  [0.]
  [0.]]]

[[[0.]
  [0.]
  [0.]
...
  [0.]
  [0.]
  [0.]]]

[[[0.]
  [0.]
  [0.]
...
  [0.]
  [0.]
  [0.]]]
=====

```

```
[21]: adv.reshape(-1,adv.shape[2]).shape
```

```
[21]: (3799, 2948)
```

```
[22]: FPredAdv = F.predict(adv.reshape(-1,adv.shape[2]))
print("=====\n\nF Predict ADV\n\n")
print(FPredAdv)
print("=====\n\n")
```

```
=====
```

```
F Predict ADV
```

```
[0.0000000e+00]
[0.0000000e+00]
[0.0000000e+00]
...
[1.3551286e-06]
[0.0000000e+00]
[0.0000000e+00]
```

```
=====
```

```
[23]: adv_pdt = np.argmax(FPredAdv, axis=1)
print("=====\n\nF Predict ADV FUNC- ADV_PDT\n\n")
print(adv_pdt)
print("=====\n\n")
print(np.unique(adv_pdt,return_counts=True))
```

```
=====
```

```
F Predict ADV FUNC- ADV_PDT
```

```
[0 0 0 ... 0 0 0]
```

```
=====
```

```
(array([0], dtype=int64), array([3799], dtype=int64))
```

```
[24]: FPredClean = F.predict(clean.reshape(-1,adv.shape[2]))
print("=====\n\nF Predict CLEAN\n\n")
print(FPredClean)
print("=====\n\n")
clean_pdt = np.argmax(FPredClean, axis=1)
```

```
print("=====\n\nF Predict CLEAN FUNC- ADV_PDT\n\n")
print(clean_pdt)
print("=====\n\n")
print(np.unique(clean_pdt,return_counts=True))
```

=====

F Predict CLEAN

```
[0.98784703]
[0.9990308 ]
[0.79184556]
...
[0.9997898 ]
[0.987847  ]
[0.9801514 ]
```

=====

=====

F Predict CLEAN FUNC- ADV\_PDT

```
[0 0 0 ... 0 0 0]
```

=====

```
(array([0], dtype=int64), array([3799], dtype=int64))
```

```
[25]: FPredPur = F.predict(purified.reshape(-1,adv.shape[2]))
print("=====\n\nF Predict PURIFIED \n\n")
print(FPredPur)
print("=====\n\n")
```

=====

F Predict PURIFIED

```
[0.05129439]
[0.05129439]
[0.05129439]
...
[0.05129439]
[0.05129439]
[0.05129439]
```



```
=====
```

```
[26]: purified_pdt = np.argmax(FPredPur, axis=1)
print("=====\n\nF Predict PURIFIED FUNC - PURIFIED_PDT\n\n")
print(purified_pdt)
print("=====\n\n")
print(np.unique(purified_pdt,return_counts=True))
```

```
=====
```

```
F Predict PURIFIED FUNC - PURIFIED_PDT
```

```
[0 0 0 ... 0 0 0]
```

```
=====
```

```
(array([0], dtype=int64), array([3799], dtype=int64))
```

```
[27]: print("=====\n\nLABEL\n\n")
print(label)
print("=====\n\n")
print(np.unique(label,return_counts=True))
```

```
=====
```

```
LABEL
```

```
[1 1 1 ... 1 1 1]
```

```
=====
```

```
(array([1]), array([3799], dtype=int64))
```

```
[28]: print(' adv acc: {:.10f},\n rct acc: {:.10f},\n\n SIMILARITY: {:.10f}'.
      ↪format( np.mean(adv_pdt==label),
              np.mean(purified_pdt==label), np.
      ↪mean(adv_pdt==purified_pdt)))
```

```
adv acc: 0.0000000000,
```

```
rct acc: 0.0000000000,
```

```
SIMILARITY: 1.0000000000
```

```
[29]: F.evaluate(clean.reshape(-1,2948), label)
```

```
3799/3799 [=====] - 2s 520us/step
```

```
[29]: [0.04282340146973827, 1.0]
```

```
[30]: F.evaluate(adv.reshape(-1,2948), label)
```

```
3799/3799 [=====] - 2s 443us/step
```

```
[30]: [299.8835265824091, 0.0]
```

```
[31]: F.evaluate(purified.reshape(-1,2948), label)
```

```
3799/3799 [=====] - 2s 447us/step
```


```
[31]: [2.970173834310955, 0.0]
```

```
[32]: FPredAdv.shape
```

```
[32]: (3799, 1)
```

```
[33]: import matplotlib.pyplot as plt  
from scipy import stats
```

```
[34]: plt.hist(FPredAdv,bins=2,color='white', edgecolor='red')  
plt.show()
```



```
[35]: StatsDesc = ['nobs', 'min,max', 'mean', 'var', 'skewness', 'kurtosis']
for i in range(len(stats.describe(FPredAdv))):
    print("=====")
    print(StatsDesc[i])
    print(stats.describe(FPredAdv)[i])
```

```
=====
nobs
3799
=====
min,max
(array([0.], dtype=float32), array([0.4306298], dtype=float32))
=====
mean
[0.00035468]
=====
var
[0.00013957]
=====
skewness
[35.453857]
=====
kurtosis
[1260.1925]
```

```
[36]: print(np.unique(purified, return_counts=True))
print(purified.shape)
c = ccc = 0
for i in range(len(purified)):
    for j in range(purified.shape[2]):
        if purified[i][0][j][0]==0:
            c+=1
        else:
            ccc+=1


print(c, ccc)
```

```
(array([0.], dtype=float32), array([11199452], dtype=int64))
(3799, 1, 2948, 1)
11199452 0
```

```
[37]: ccc/15621
```

```
[37]: 0.0
```

```
[38]: plt.hist(FPredPur,bins=2,color='white', edgecolor='red')
plt.show()
```



```
[39]: StatsDesc = ['nobs','min,max','mean','var','skewness','kurtosis']
for i in range(len(stats.describe(FPredPur))):
    print("=====")
    print(StatsDesc[i])
    print(stats.describe(FPredPur)[i])
```

```
=====
nobs
3799
=====
min,max
(array([0.05129439], dtype=float32), array([0.05129439], dtype=float32))
=====
mean
[0.05129439]
=====
var
[1.0961918e-20]
=====
skewness
[35.585575]
=====
```

```
kurtosis  
[1263.3259]
```

## 6 Conclusions FGSM PURE MAL

Given that FGSM Pure MAL is also converting all to 0. We were stuck here for quite some time and then began going over to other methods. The next section discusses about those methods

[ ]: